

Bertelsmann Tech Scholarship - Data Track

27.4 Text + Quiz: Your First JOIN

A query using a JOIN statement.

- SELECT clause indicates column(s) of data will see in the output
- FROM clause indicates 1st table we're pulling data.
- JOIN indicates 2nd table. JOINing the matching PK-FK links from 2 tables.
- ON clause specifies col on which you'd like to merge 2 tables together.

-->>> Notice > we pull data from 2 tables: orders - accounts

but SELECT only pulling data from orders table, reference columns from orders table.

ON statement holds 2 columns that get linked across 2 tables, will be the focus in the next concepts.

--->>> If we wanted to only pull individual elements from either the orders or accounts table, using the exact same information in the FROM and ON statements. However, SELECT, will need to know how to specify tables and columns in the SELECT:

This query only pulls two columns

```
SELECT accounts.name,  
       orders.occurred_at
```

```
FROM orders
```

```
JOIN accounts
```

```
ON orders.account_id = accounts.id;
```

```
SELECT orders.*
```

```
FROM orders
```

```
JOIN accounts
```

```
ON orders.account_id = accounts.id;
```

table name column

This query pulls all columns from both accounts + orders table.

```
SELECT *
```

```
FROM orders
```

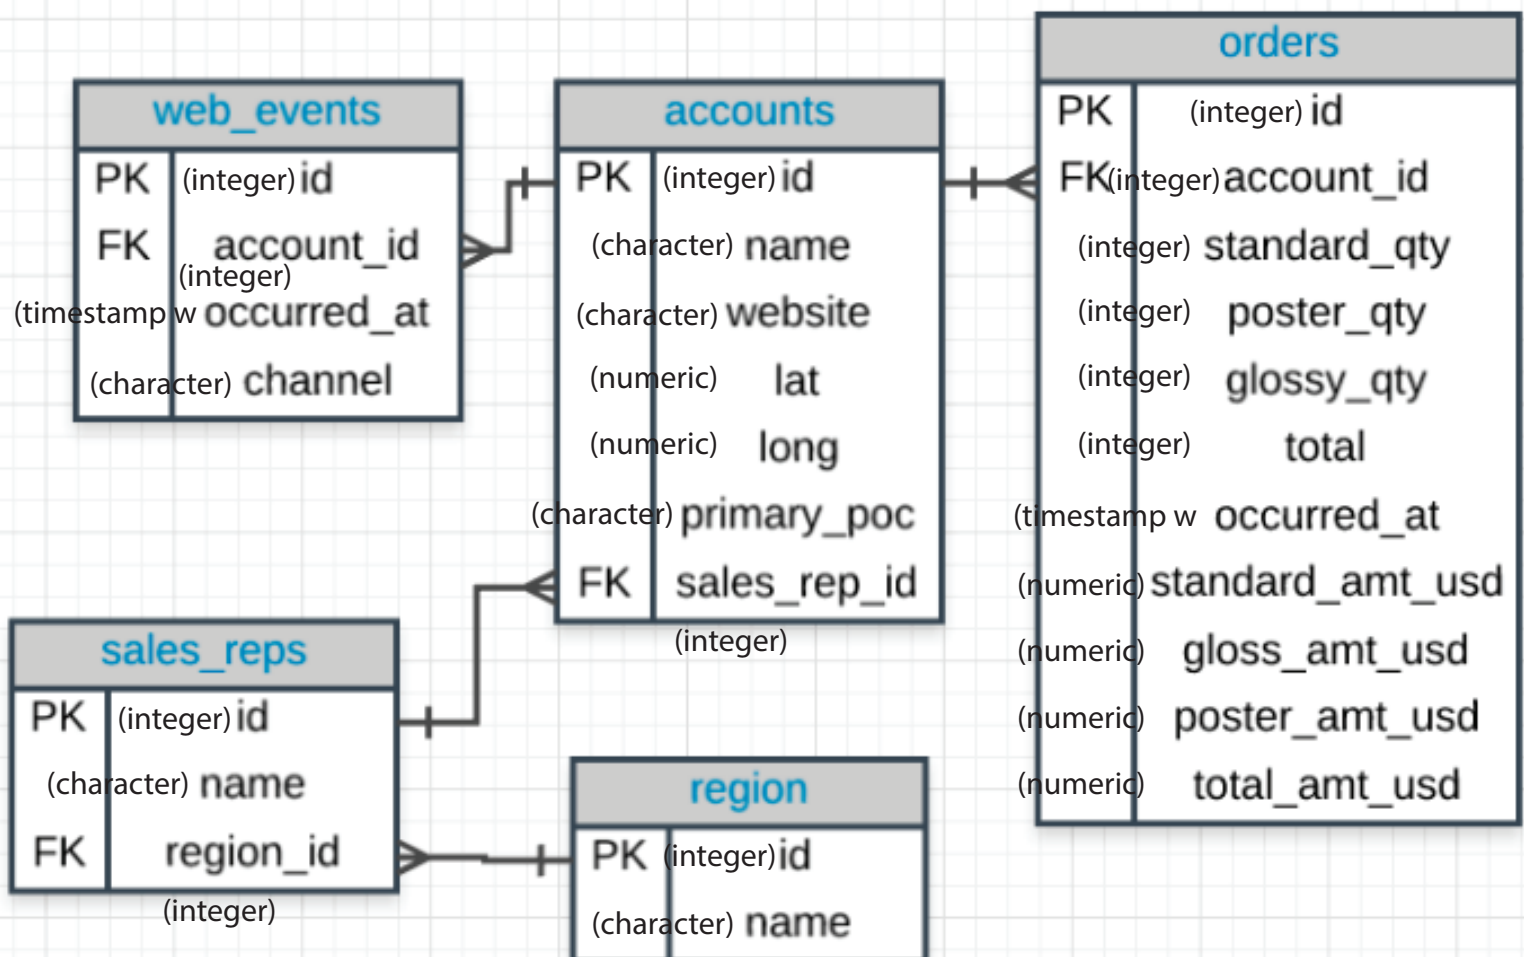
```
JOIN accounts
```

```
ON orders.account_id = accounts.id;
```

27.6 Text: ERD Reminder = Entity Relationship Diagrams

PK = primary key = a col has a unique value for every row ~ id PK is usually the first column in our table.

Foreign Key - FK = a column in one table that is a primary key in a different table >> region_id
the foreign key is when the primary key appears in another table, account_id
but it doesn't need to be unique. sales_rep_id



Parch & Posey database

Bertelsmann Tech Scholarship - Data Track

27.7 Text: Primary PK and Foreign Keys FK

PK = primary key = a col has a unique value for every row ~ id, common 1st col in each of our tables, in most dbs.

FK = a column in one table that is a primary key in a different table.

FK linked to PK of another table.

27.9 Text + Quiz: JOIN Revisited

```
SELECT orders.*
```

```
FROM orders
```

```
JOIN accounts
```

```
ON orders.account_id = accounts.id;
```

- SQL query want to join 2 table - 1 in FROM and the other in JOIN
Then in ON, ALWAYS have PK = FK;

- The actual ordering of which table name goes first in this statement doesn't matter so much. So, we could also write

```
ON accounts.id = orders.account_id;
```

- JOIN More than Two Tables

This same logic can actually assist in joining more than two tables together. To join all 3 tables, use the same logic.

```
SELECT *
```

```
FROM web_events
```

```
JOIN accounts
```

```
ON web_events.account_id = accounts.id
```

```
JOIN orders
```

```
ON accounts.id = orders.account_id
```

- To pull specific cols, SELECT need to specify table + col name. Pull only 3 cols in these 3 tables by changing SELECT, but maintain rest of JOIN info

```
SELECT web_events.channel, accounts.name, orders.total
```

27.10 Video: Alias

- When JOIN tables, nice to give table an alias = 1st letter of table name.

```
SELECT orders.*
```

```
FROM orders
```

```
JOIN accounts
```

```
ON orders.account_id = accounts.id;
```

```
SELECT o.*, a.*
```

```
FROM orders o
```

```
JOIN accounts a
```

```
ON o.account_id = a.id;
```

27.11 Quiz: JOIN Questions Part I

```
SELECT a.primary_poc, w.occurred_at, w.channel, a.name
```

```
FROM web_events w
```

```
JOIN accounts a
```

```
ON w.account_id = a.id
```

```
WHERE a.name = 'Walmart';
```

```
SELECT r.name region, s.name rep, a.name account
```

```
FROM sales_reps s
```

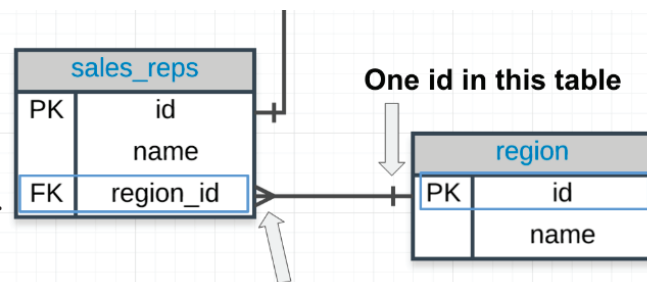
```
JOIN region r
```

```
ON s.region_id = r.id
```

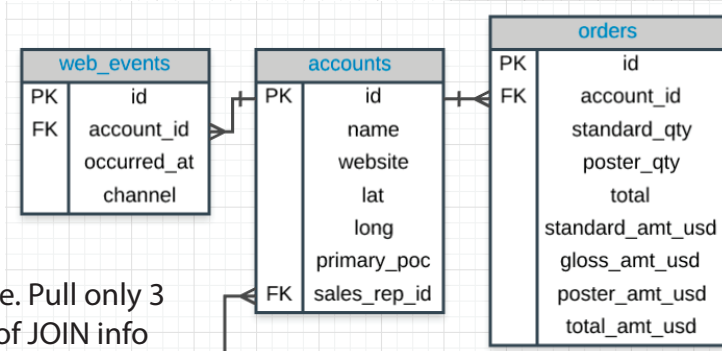
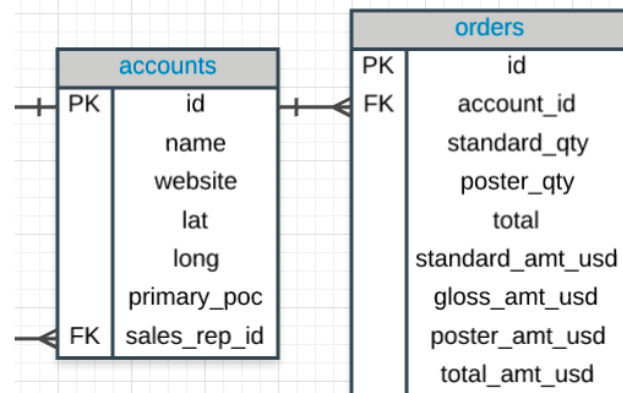
```
JOIN accounts a
```

```
ON a.sales_rep_id = s.id
```

```
ORDER BY a.name;
```



Many of this id in this table
primary-foreign key link that connects 2 tables.
crow's foot shows FK can appear in many rows in table
PK single line shows id appears only once per row in table



```
SELECT r.name region, a.name account,
       o.total_amt_usd/(o.total + 0.01) unit_price
FROM region r
JOIN sales_reps s
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
JOIN orders o
ON o.account_id = a.id;
```


27.13 Video: Motivation for Other JOINS

INNER JOIN = 1-to-1 and 1-to-many relationships when we introduced PKs and FKs. Notice, traditional dbs do not allow for many-to-many relationships, as these break the schema down pretty quickly.

The types of relationships that exist in a db matter less to analysts, but you do need to understand why you would perform diff types of JOINS, & what data you are pulling from db. These ideas 'll be expanded upon in next concepts.

27.14 LEFT and RIGHT JOINS

OUTER JOIN = JOIN data depending on the question we are asking. Notice each of these new JOIN states pulls all the same rows as an INNER JOIN, which you saw by just use JOIN, but they also potentially pull some additional rows. If there is not matching info in JOINed table, then will have cols w/ empty cells, without data = introduce a new data type = NULL.



27.15 Text: Other JOIN Notes

INNER JOIN = pulled rows only if they exist as a match across 2 tables.

OUTER JOIN = allow us to pull rows that might only exist in 1 of the 2 tables, into a new data type NULL.

LEFT OUTER JOIN = LEFT JOIN = obtain all rows of INNER JOIN, & additional rows from the table in the FROM.

The last type of join = full outer join. This will return the inner join result set, as well as any unmatched rows from either of 2 tables being joined. FULL OUTER JOIN, same as OUTER JOIN.

examples of outer join =

<https://www.w3resource.com/sql/joins/perform-a-full-outer-join.php>

If we were to flip the tables, we would get same exact result as the JOIN statement:

27.18 Video: JOINS and Filtering

- When db executes query, it executes JOIN and ON clause first

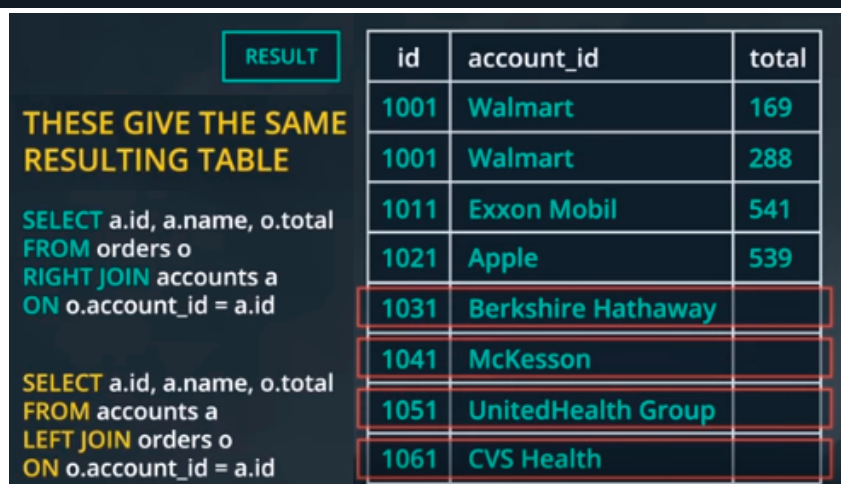
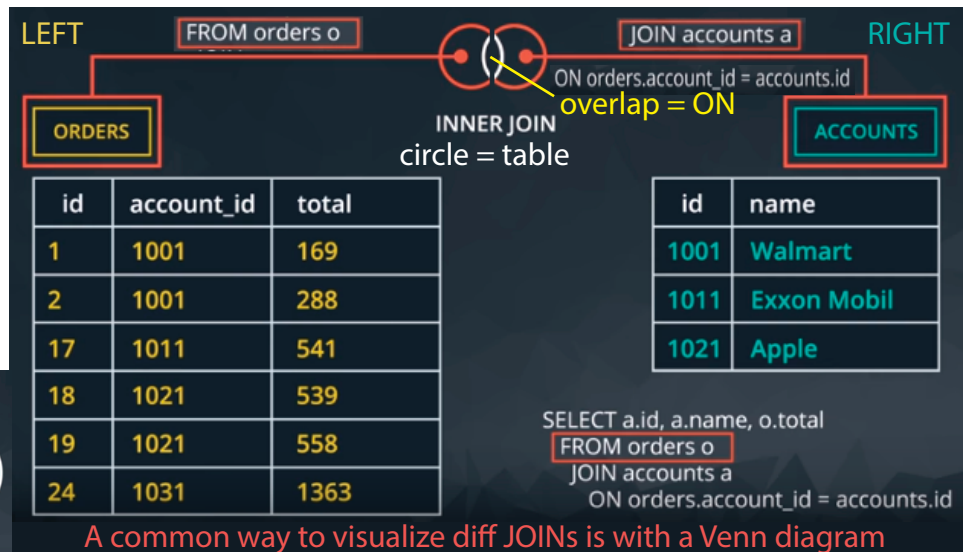

=> building new result set

=> then filtered using the WHERE clause.

- ON clause of an INNER JOIN will produce the same result as keeping it in the WHERE clause.

**LOGIC IN THE ON CLAUSE
REDUCES THE ROWS BEFORE
COMBINING THE TABLES**

**LOGIC IN THE WHERE CLAUSE
OCCURS AFTER THE
JOIN OCCURS**



5	Sri Lanka	NULL
6	Brazil	NULL

```
SELECT orders.*,
       accounts.*
FROM demo.orders
LEFT JOIN demo.accounts
ON orders.account_id = accounts.id
WHERE accounts.sales_rep_id = 321500
```

AND replace WHERE, as the continue of ON clause, will add more rows.

Bertelsmann Tech Scholarship - Data Track

27.19 Quiz: Last Check 27.20 Solutions: Last Check

```
SELECT r.name region, s.name rep, a.name account
FROM sales_reps s
JOIN region r
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
WHERE r.name = 'Midwest'
ORDER BY a.name;
```

SOLUTION 1

```
SELECT r.name region, s.name rep, a.name account
FROM sales_reps s
JOIN region r
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
WHERE r.name = 'Midwest' AND s.name LIKE 'S%'
ORDER BY a.name;
```

SOLUTION 2

```
SELECT r.name region, s.name rep, a.name account
FROM sales_reps s
JOIN region r
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
WHERE r.name = 'Midwest' AND s.name LIKE 'K%'
ORDER BY a.name;
```

SOLUTION 3

```
SELECT DISTINCT o.occurred_at, a.name, o.total, o.total_amt_usd
FROM accounts a
JOIN orders o
ON a.id = o.account_id
WHERE o.occurred_at BETWEEN '01-01-2015' AND '01-01-2016'
ORDER BY o.occurred_at DESC;
```

SOLUTION 8

27.21 Text: Recap & Looking Ahead

PK - are unique for every row in a table, generally 1st col in db (like id col for every table in our db).

FK - are primary key appearing in another table, which allows the rows to be non-unique.

Choosing the set up of data in our database is very important, but not usually the job of a data analyst. This process is known as Database Normalization.

JOINS = combine data from multiple tables, 3 JOIN statements most likely to use are:

- JOIN - an INNER JOIN that only pulls data that exists in both tables.

- LEFT JOIN - pulls all data exists in both tables, as well as all of the rows from table in FROM even if they do not exist in the JOIN statement.

- RIGHT JOIN - pulls all data exists in both tables, as well as all of the rows from the table in the JOIN even if they do not exist in the FROM statement.

- Advanced JOINS used in very specific use cases. UNION and UNION ALL, CROSS JOIN, and the tricky SELF JOIN.

Alias tables and cols using AS or not using it, allows to be more efficient in the number of characters you need to write, while at the same time you can assure that your column headings are informative of the data in your table.

- Next is aggregating data. Our SQL might still a bit disconnected from statistics and using Excel like platforms.

Aggregations allow to write SQL code more complex queries, which assist in answering questions like:

Which channel generated more revenue?

Which account had an order with the most items?

Which sales_rep had the most orders? or least orders? How many orders did they have?

```
SELECT r.name region, a.name account,
       o.total_amt_usd/(o.total + 0.01) unit_price
FROM region r
JOIN sales_reps s
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
JOIN orders o
ON o.account_id = a.id
WHERE o.standard_qty > 100;
```

SOLUTION 4

```
SELECT r.name region, a.name account,
       o.total_amt_usd/(o.total + 0.01) unit_price
FROM region r
JOIN sales_reps s
ON s.region_id = r.id
JOIN accounts a
ON a.sales_rep_id = s.id
JOIN orders o
ON o.account_id = a.id
WHERE o.standard_qty > 100 AND o.poster_qty > 50
ORDER BY unit_price;
```

SOLUTION 5

SOLUTION 6

```
SELECT DISTINCT a.name, w.channel
FROM accounts a
JOIN web_events w
ON a.id = w.account_id
WHERE a.id = '1001';
```

SOLUTION 7

SELECT DISTINCT
to narrow down the results to
only the unique values.